

# FreeBSD Concrete Architecture

pink fluffy unicorns



# OUR TEAM

Benedict Miguel  
Jiachen Wang  
Keshav Gainda  
Ketan Bhandari  
Mohaimen Hassan



Pegah Fallah  
Shami Sharma  
Sidharth Sudarsan  
Suryam Thaker  
Xiaochuan Wang



# AGENDA

## 01 Derivation Process

Subsystem assignments and derivation of contain files

## 02 Top Level Architecture

Concrete architecture, style and functionality

## 03 Interactions

Between subsystems

## 04 Reflexion Analysis

Conceptual model  
vs  
Reflexion model

## 05 LSEdit Demo

Extracted architecture and subsystems demo on LSEdit

## 06 Lessons Learned

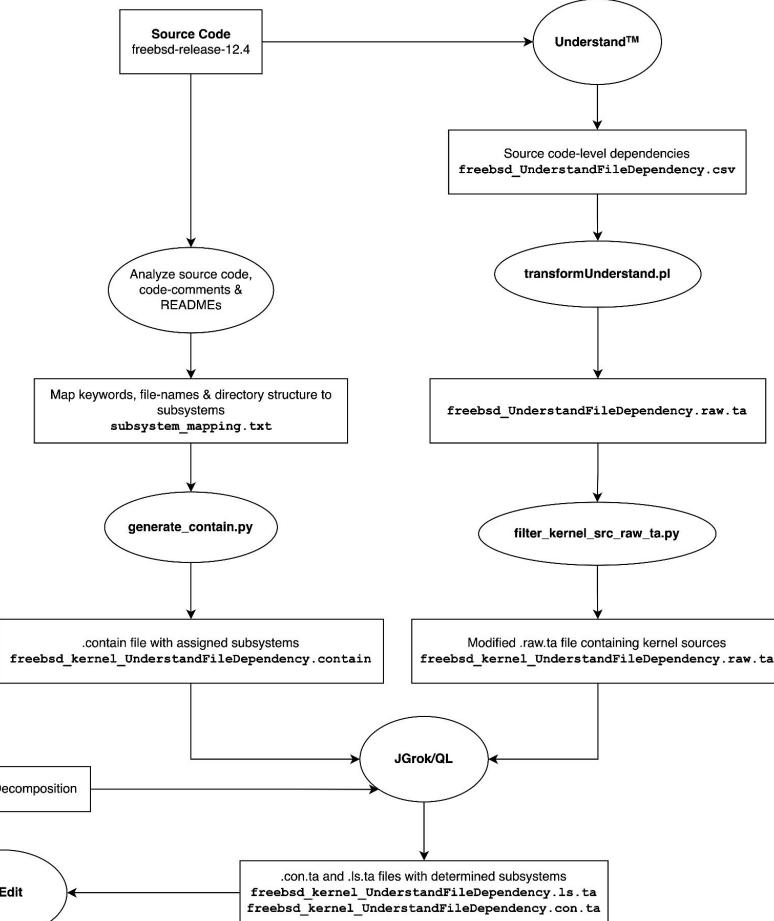
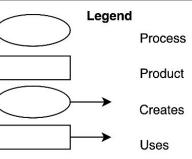


# Tools

- **Understand** by SciTools: Static Code Analysis Tool
  - Source code-level dependencies:
    - Extracts functions calls and variable access relations  
(e.g. function y calls function x)
    - Control flow and data flow dependencies
- **JGrok/QL**
  - Determine relations between files and subsystems
- **LSEdit**: The Graphical Landscape Editor
  - Visualize extracted system architecture
- **Python**
  - Scripting for processing, filtering and generating data



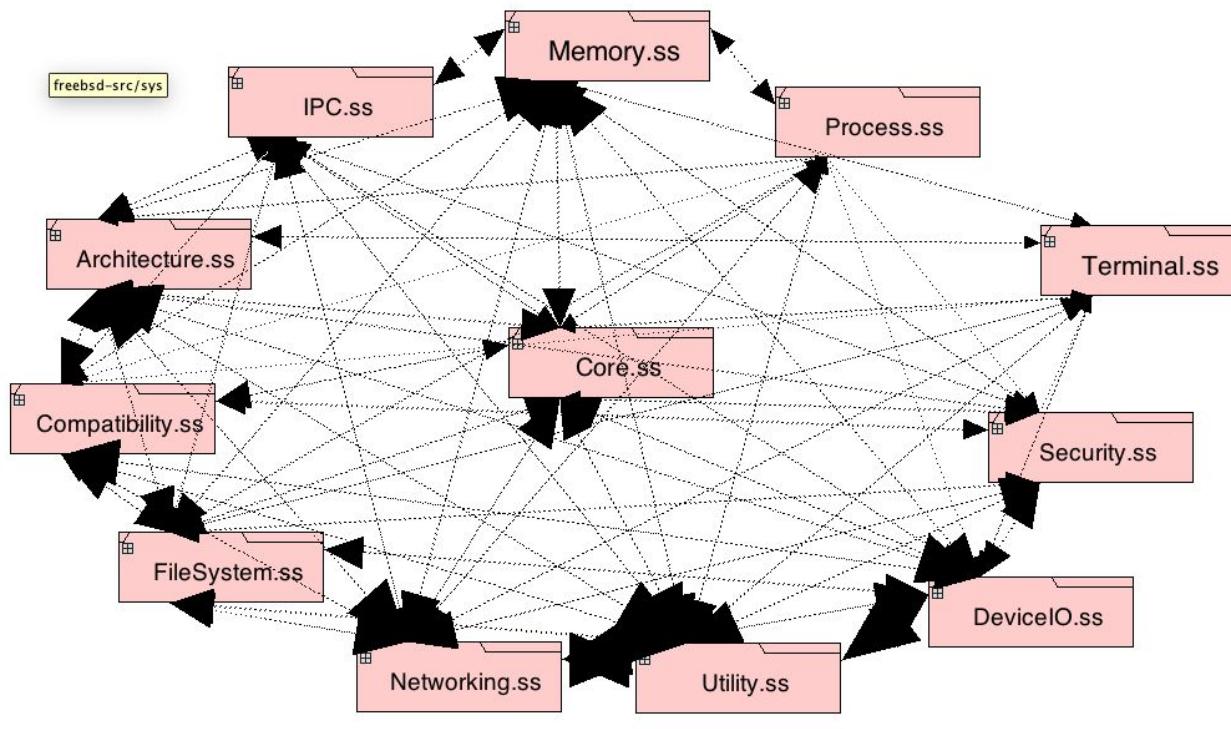
# DERIVATION PROCESS



- /sys/ contains the source files of the kernel
- Assigned subsystems based on the subfolder structure of /sys/
- Determined the subfolders that may or may not be associated with a subsystem
- Analyzed the file-naming, code comments, READMEs and source-code function headers



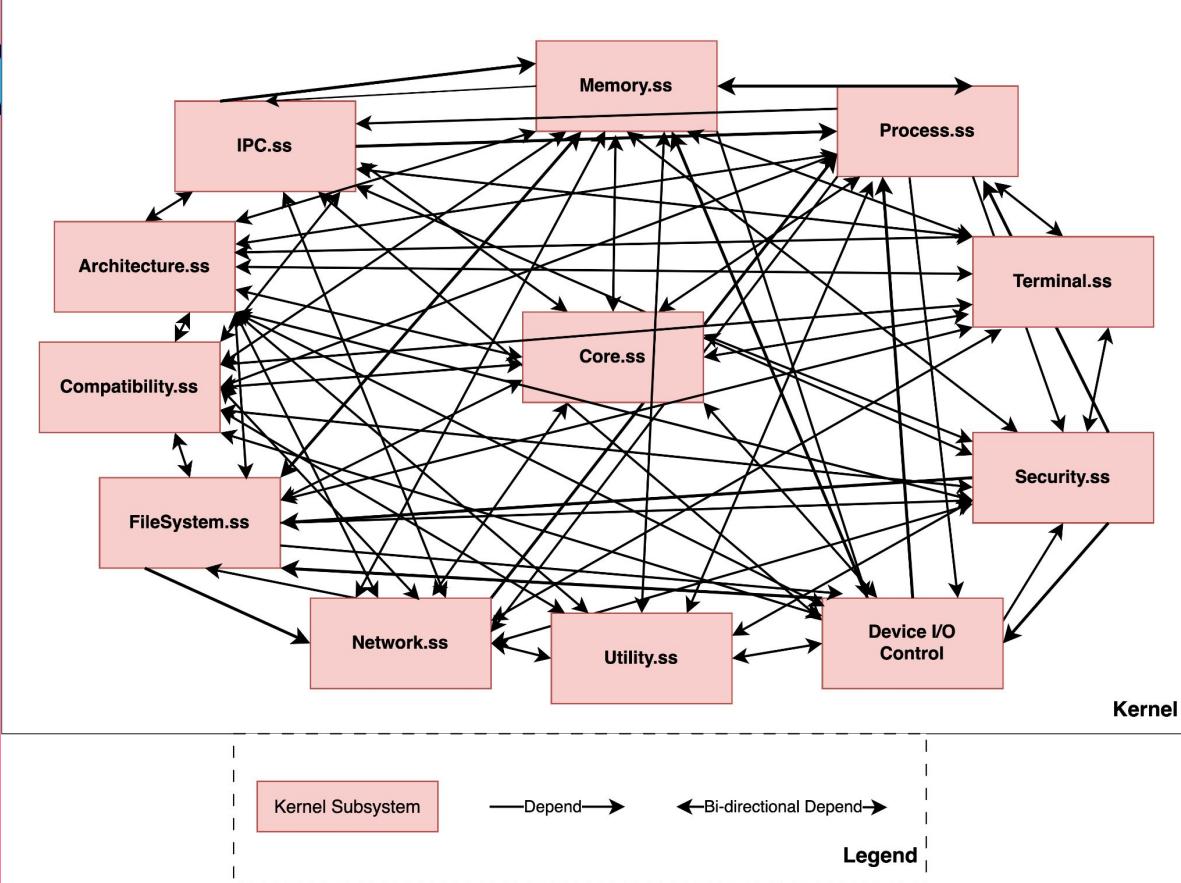
## Top-Level Concrete Architecture



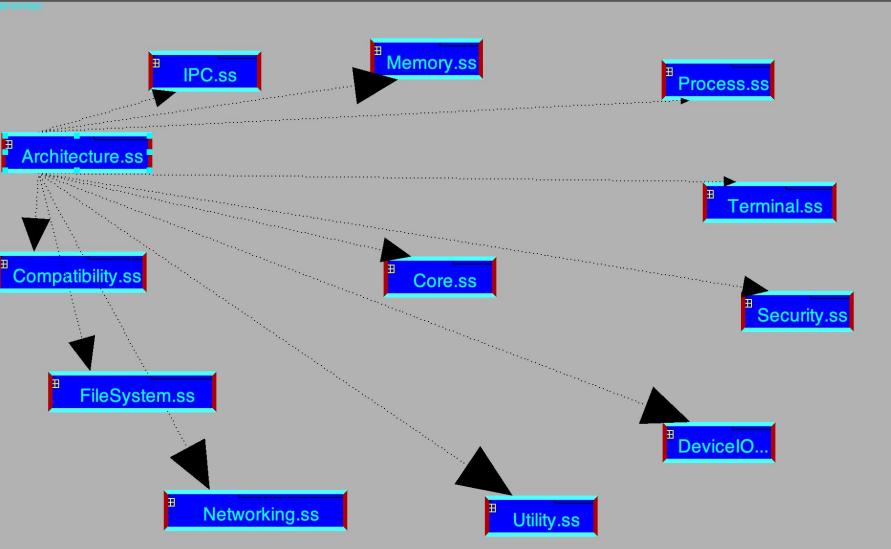
<code>freebsd-src/sys (12)</code>
Architecture.ss (2071)
Core.ss (133)
DeviceIO.oss (5562)
FileSystem.ss (276)
IPC.ss (16)
Memory.ss (61)
Networking.ss (880)
Process.ss (13)
Security.ss (71)
Terminal.ss (10)
Utility.ss (682)
Compatibility.ss (255)

# Top-Level Concrete Architecture

Monolithic Kernel

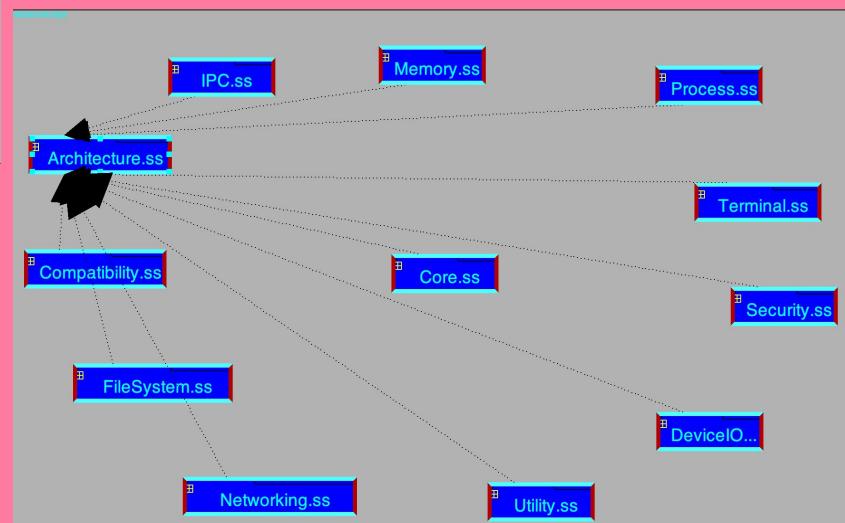


# ARCHITECTURE

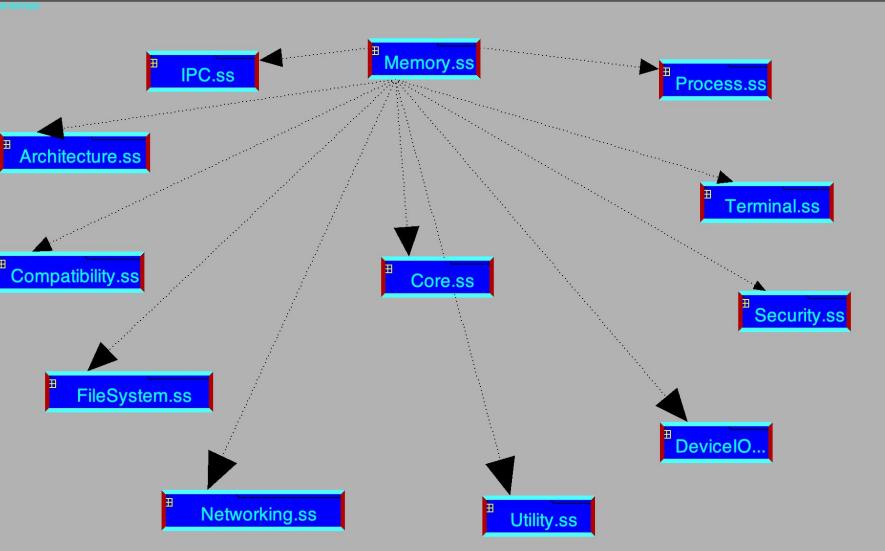


amd64      risc-v  
arm64      x86  
i386      mips  
powerpc      sparc64

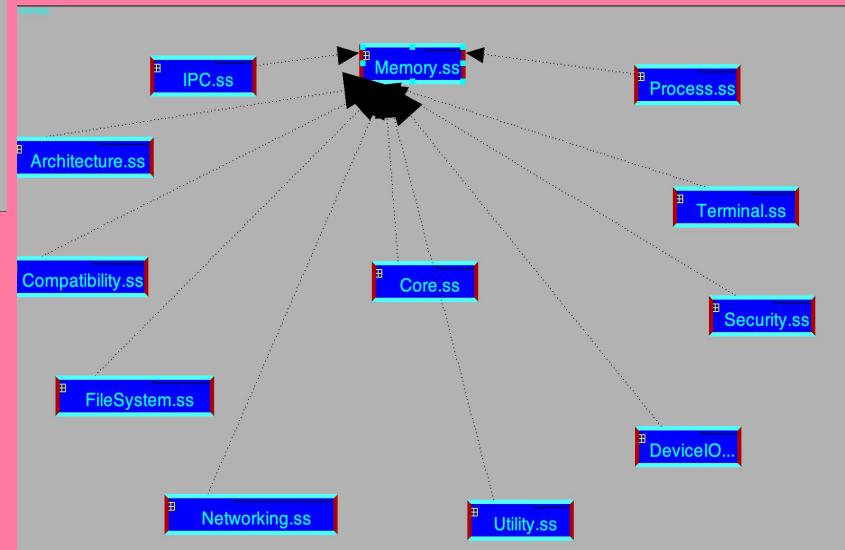
## Architecture-dependent code



# MEMORY

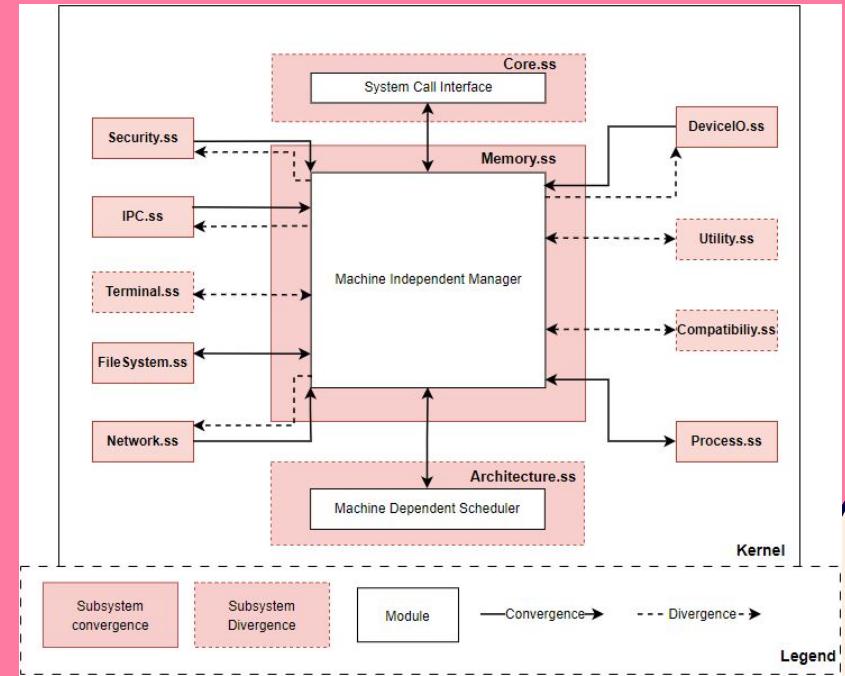
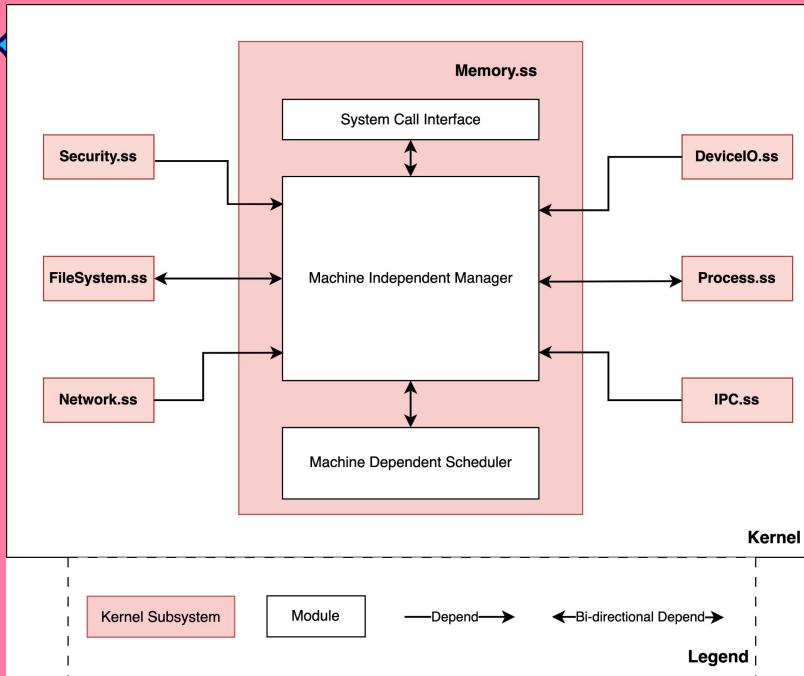


- Implementing the virtual memory abstraction
- Syscalls: *mlock, execve, read, write, etc.*
- Page-fault handling
- Kernel memory allocation
  - *malloc, free*



- Handling memory pressure
  - Reclaiming pages
  - Approximating LRUs
- Coupled with miscellaneous subsystems
  - *tmpfs, pmap*
  - Scalability w.r.t CPU & RAM

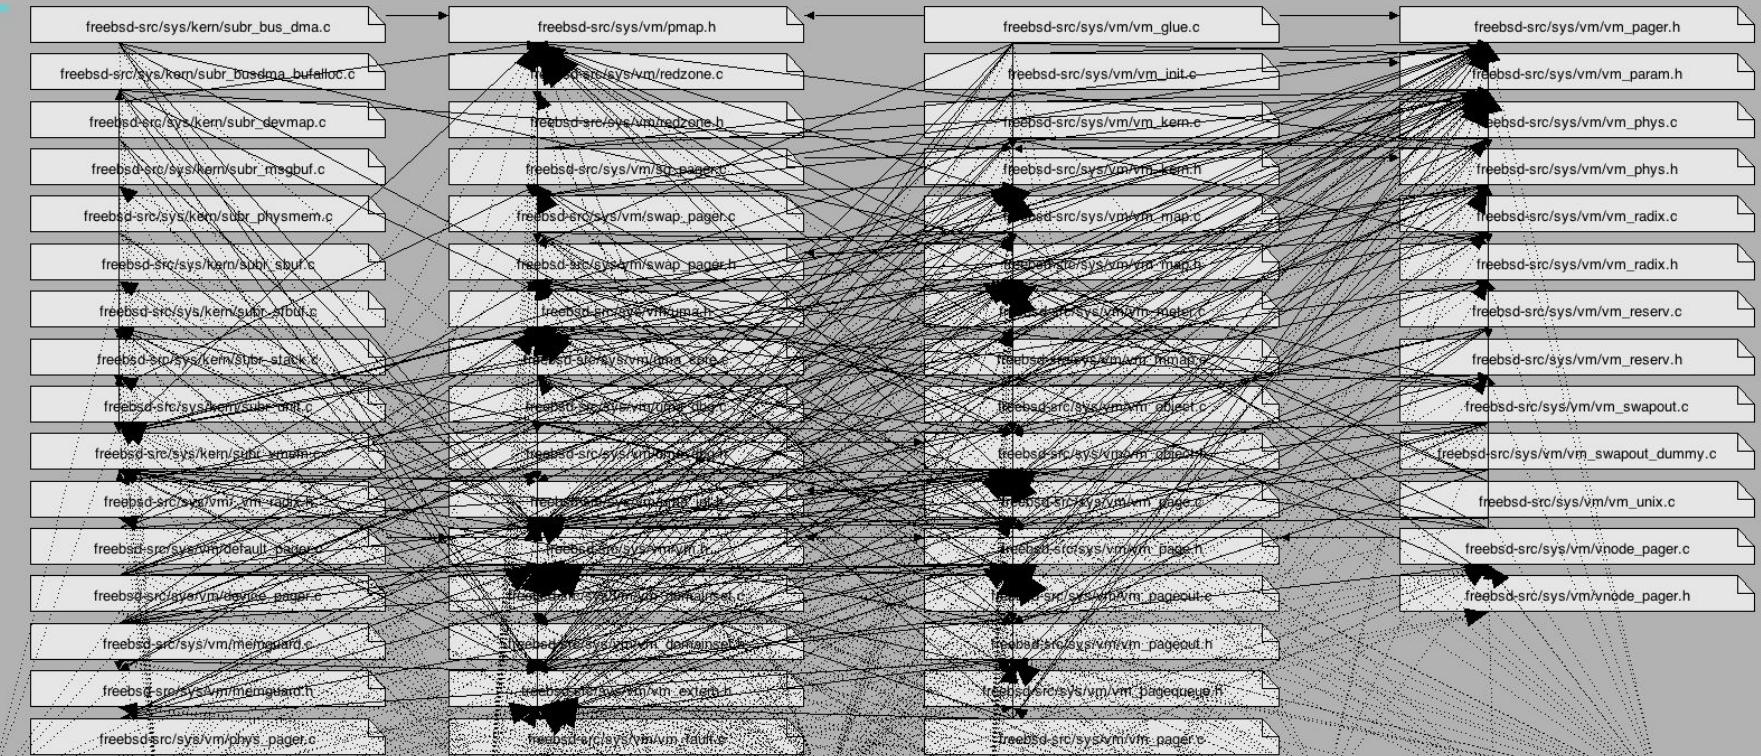
# MEMORY SUBSYSTEM



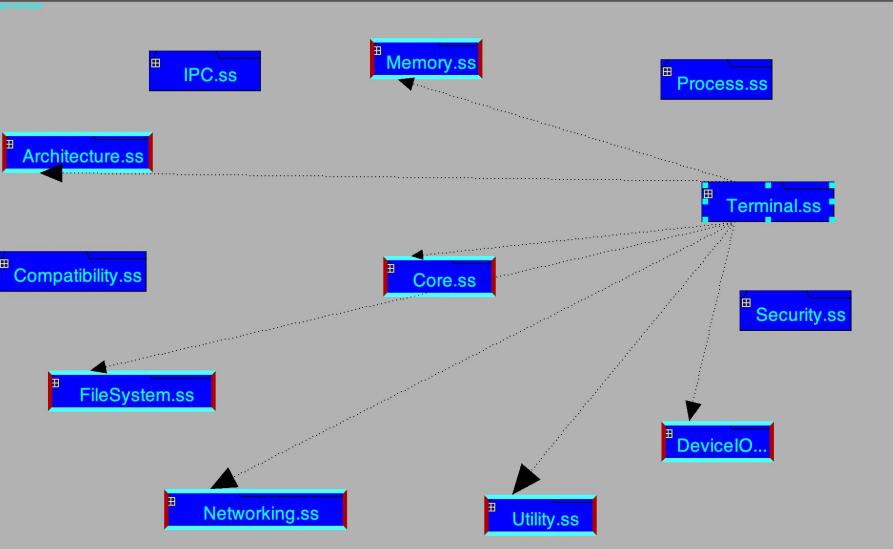


# MEMORY: LS EDIT MATRIX LAYOUT

Memory.h



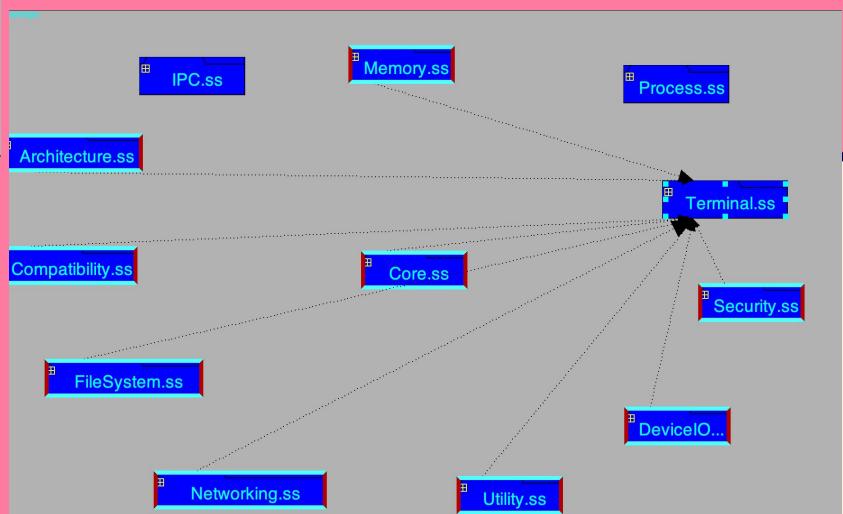
# TERMINAL



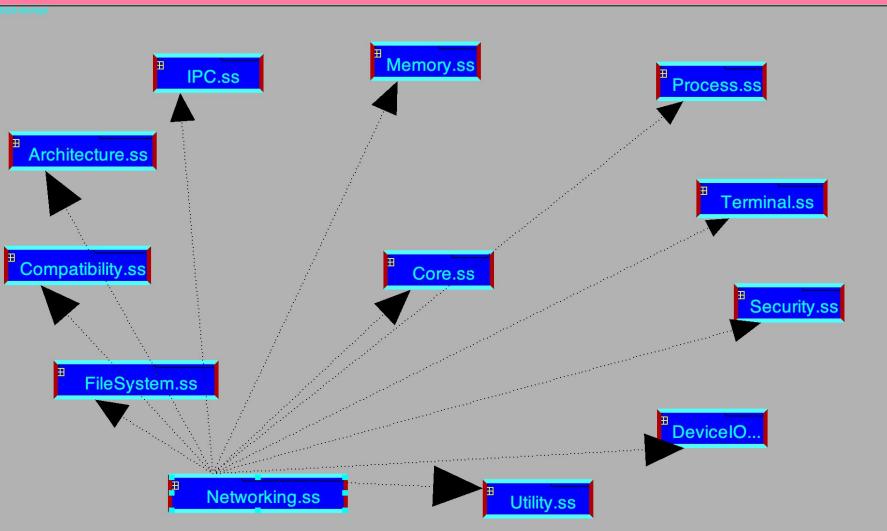
- Terminal driver
- Terminal emulator
- Virtual Terminal Manager

## User Interactions

- Manage the interactions between the user and the system through an interface



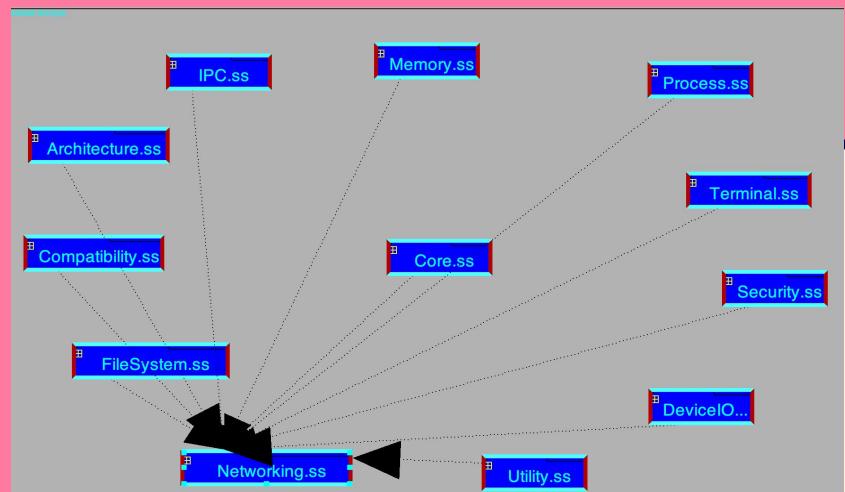
# NETWORKING



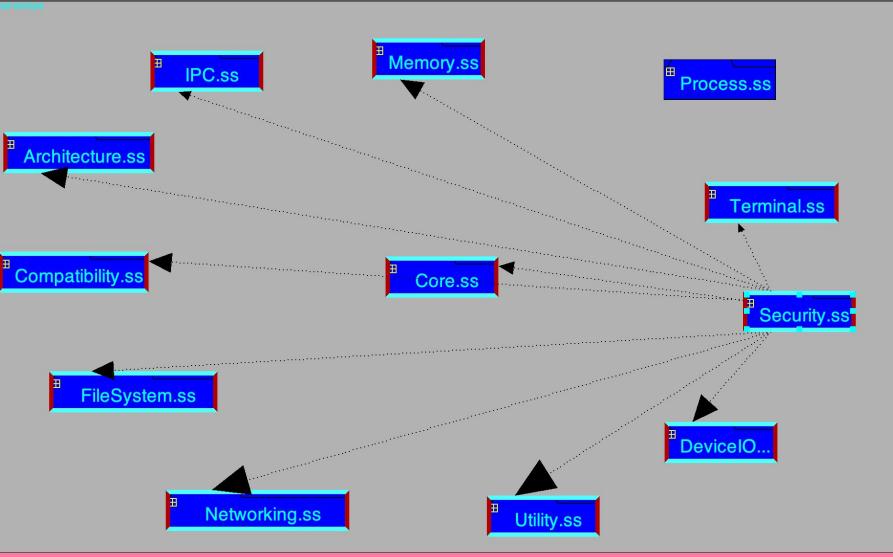
`/sys/net` - core networking code  
`/sys/net80211` - wireless networking (IEEE 802.11)  
`/sys/netgraph` - graph-based networking subsystem  
`/sys/netinet` - IPv4 protocol implementation

`/sys/netinet6` - IPv6 protocol implementation  
`/sys/netipsec` - IPsec protocol implementation  
`/sys/netpfifl` - packet filters

`/sys/net` - core networking code  
`/sys/net80211` - wireless networking (IEEE 802.11)  
`/sys/netgraph` - graph-based networking subsystem  
`/sys/netinet` - IPv4 protocol implementation



# SECURITY

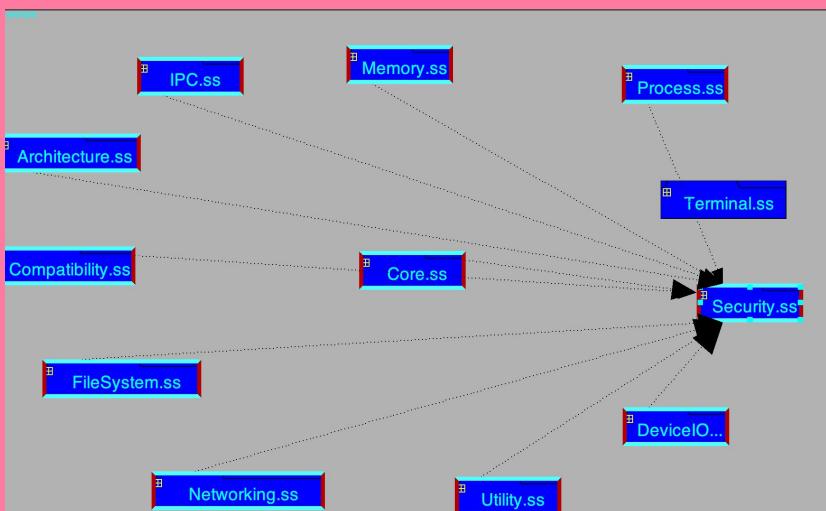


contain **Security.ss** **Cryptography.ss**

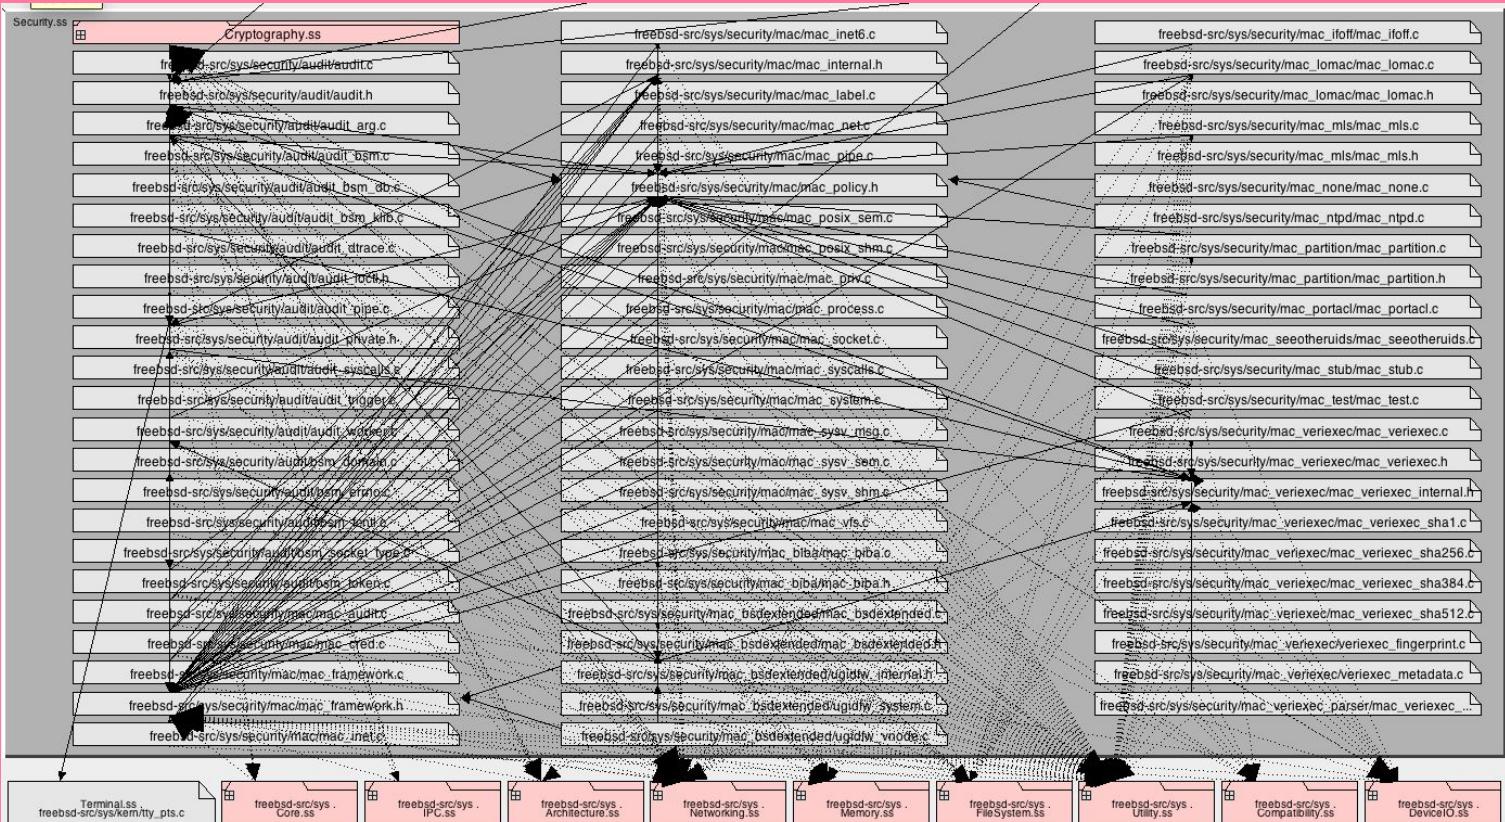
*opencrypto* is a framework for drivers of cryptographic hardware to register with the kernel so ``consumers'' (other kernel subsystems, and eventually users through an appropriate device) are able to make use of it.

Implements Multilevel Security policies:

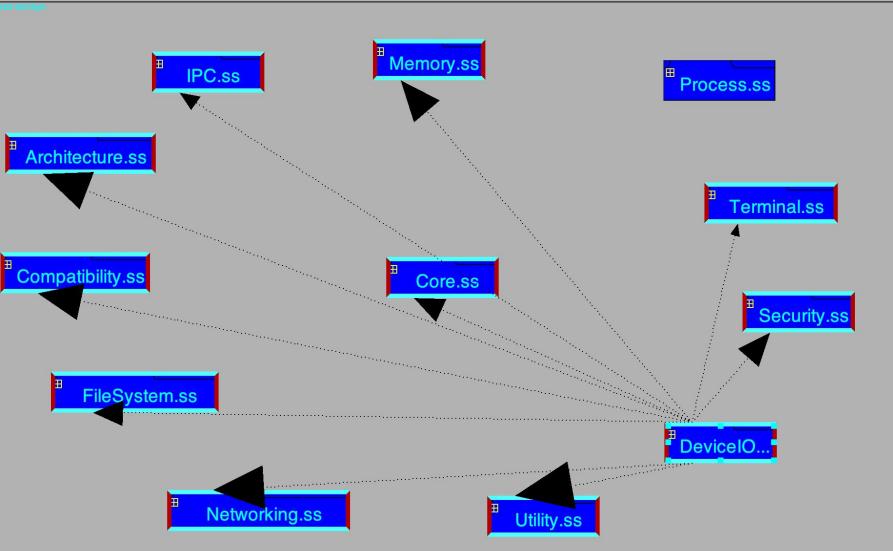
- Mandatory access control (MAC)
- Discretionary access control (DAC)



# SECURITY: LSEDIT MATRIX LAYOUT



# DEVICE IO



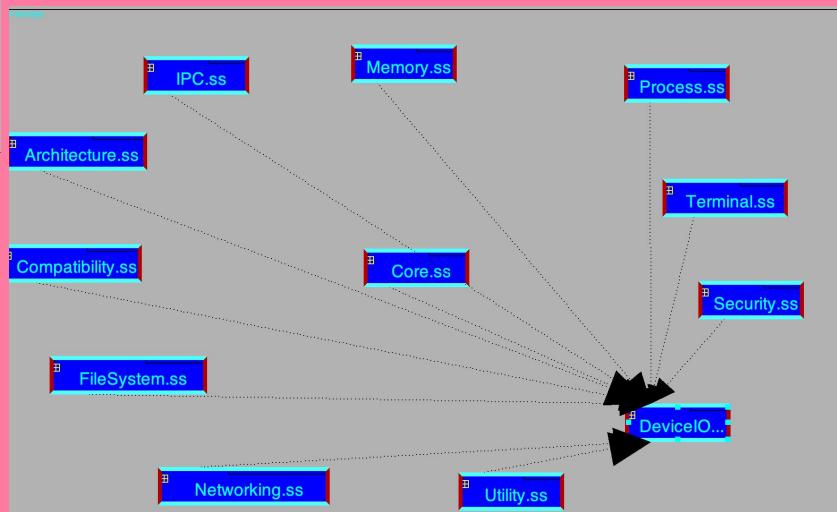
## Device driver interface

- Hardware to OS

## Bus management

- Build connections to the bus

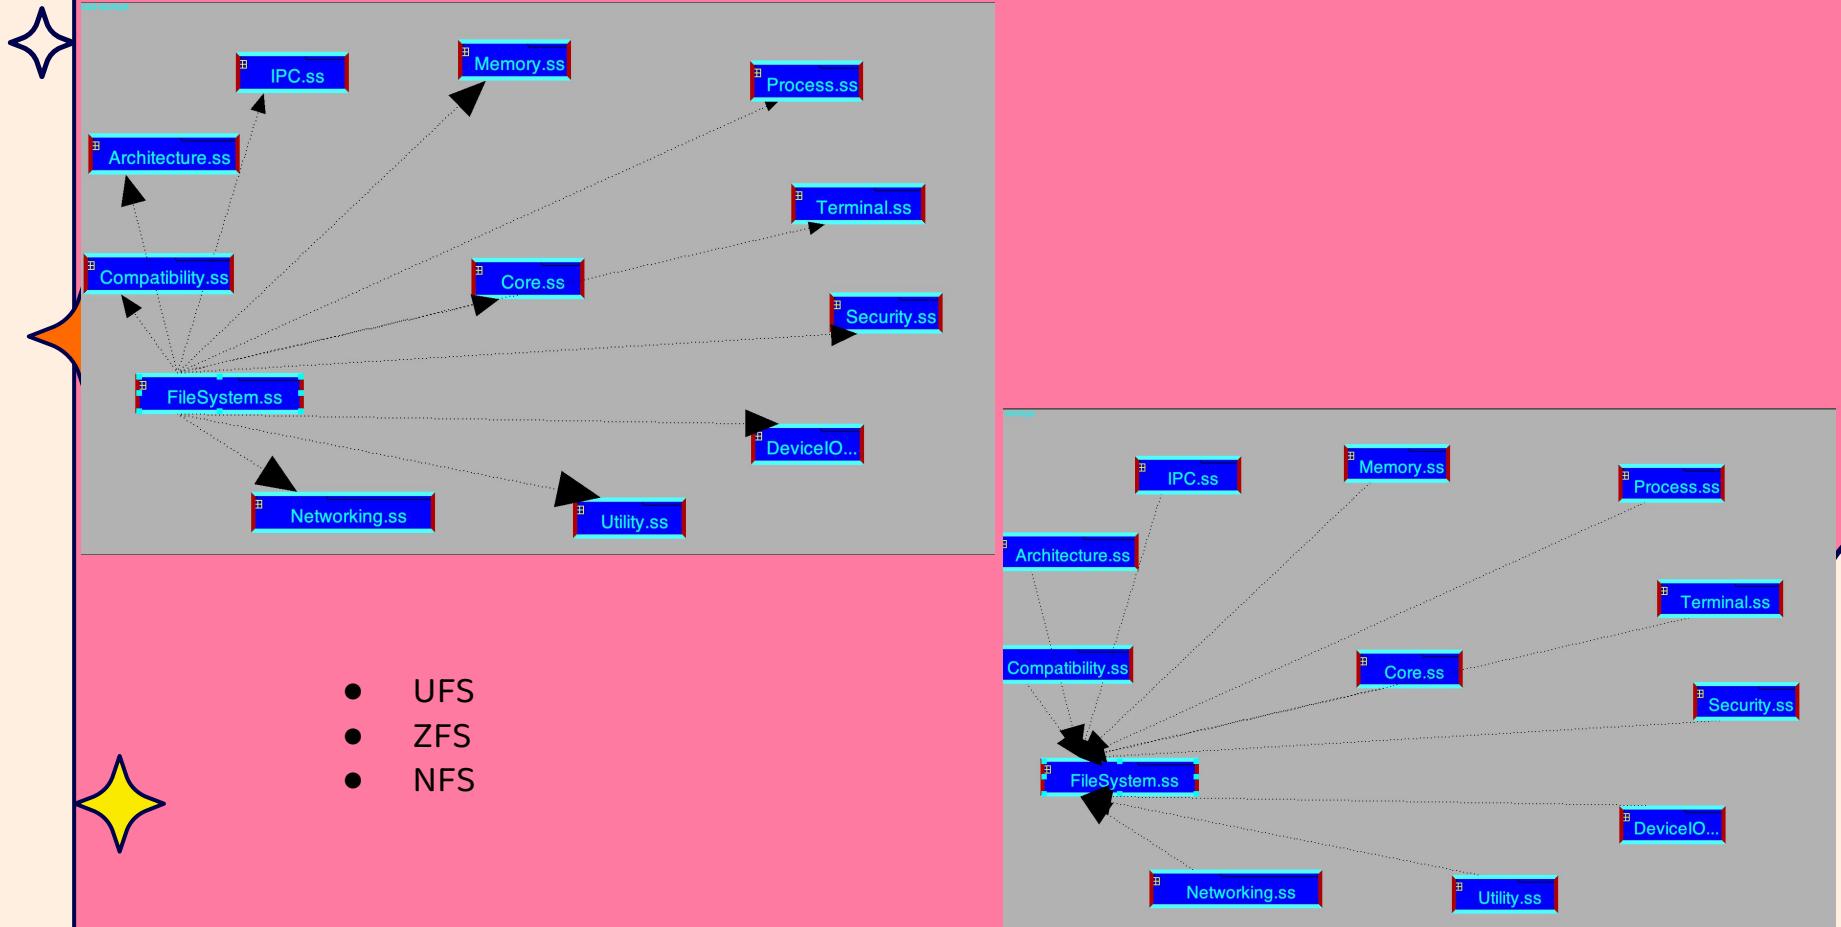
- Device allocation and management
- I/O scheduling
- Buffer management



# I/O Device Management

- Kernel acts as an intermediary between the hardware and the operating system
- Responsible for providing the abstract view of the hardware to the operating system and handling all I/O requests.
- FreeBSD provides more advanced I/O device management systems
  - Complete Fair Queueing (CFQ)**
  - Deadline Scheduling**
- Able to handle very large amounts of I/O
- Ideal in enterprise environments where the I/O requirements are very demanding

# FILE SYSTEM



# File Management

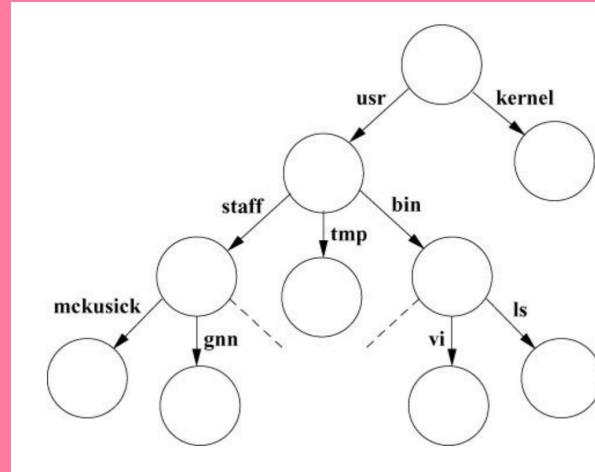
- Uses Unix File System (UFS) to store files.
- Files are managed in a tree structure
- Also provides a graphical method of file management

## Zettabyte File System

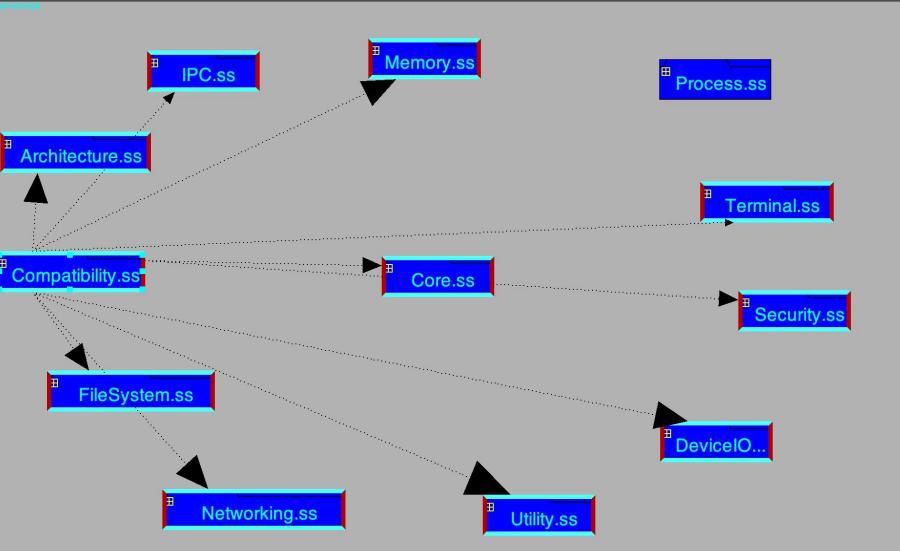
- Data Compression, Check Summing, Snapshots

## Network File System

- Access to files over the network
- Been known to support various other file-management systems.
- Very ideal in multi-device environments.

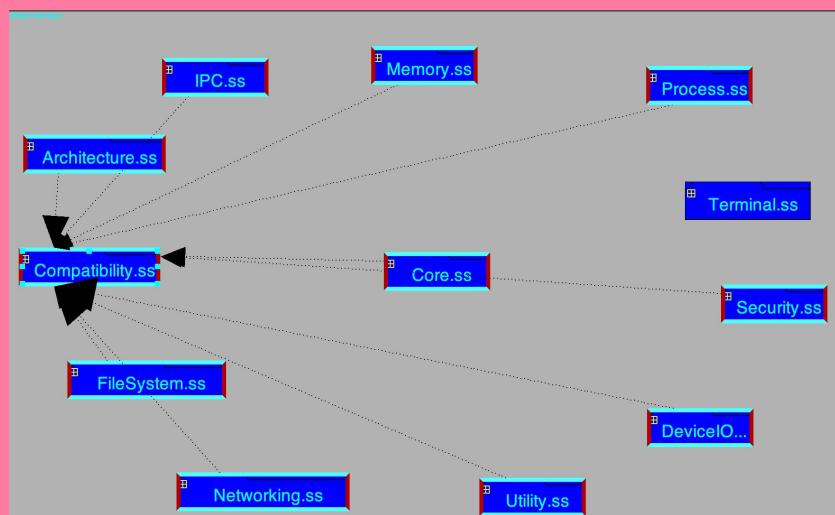


# COMPATIBILITY

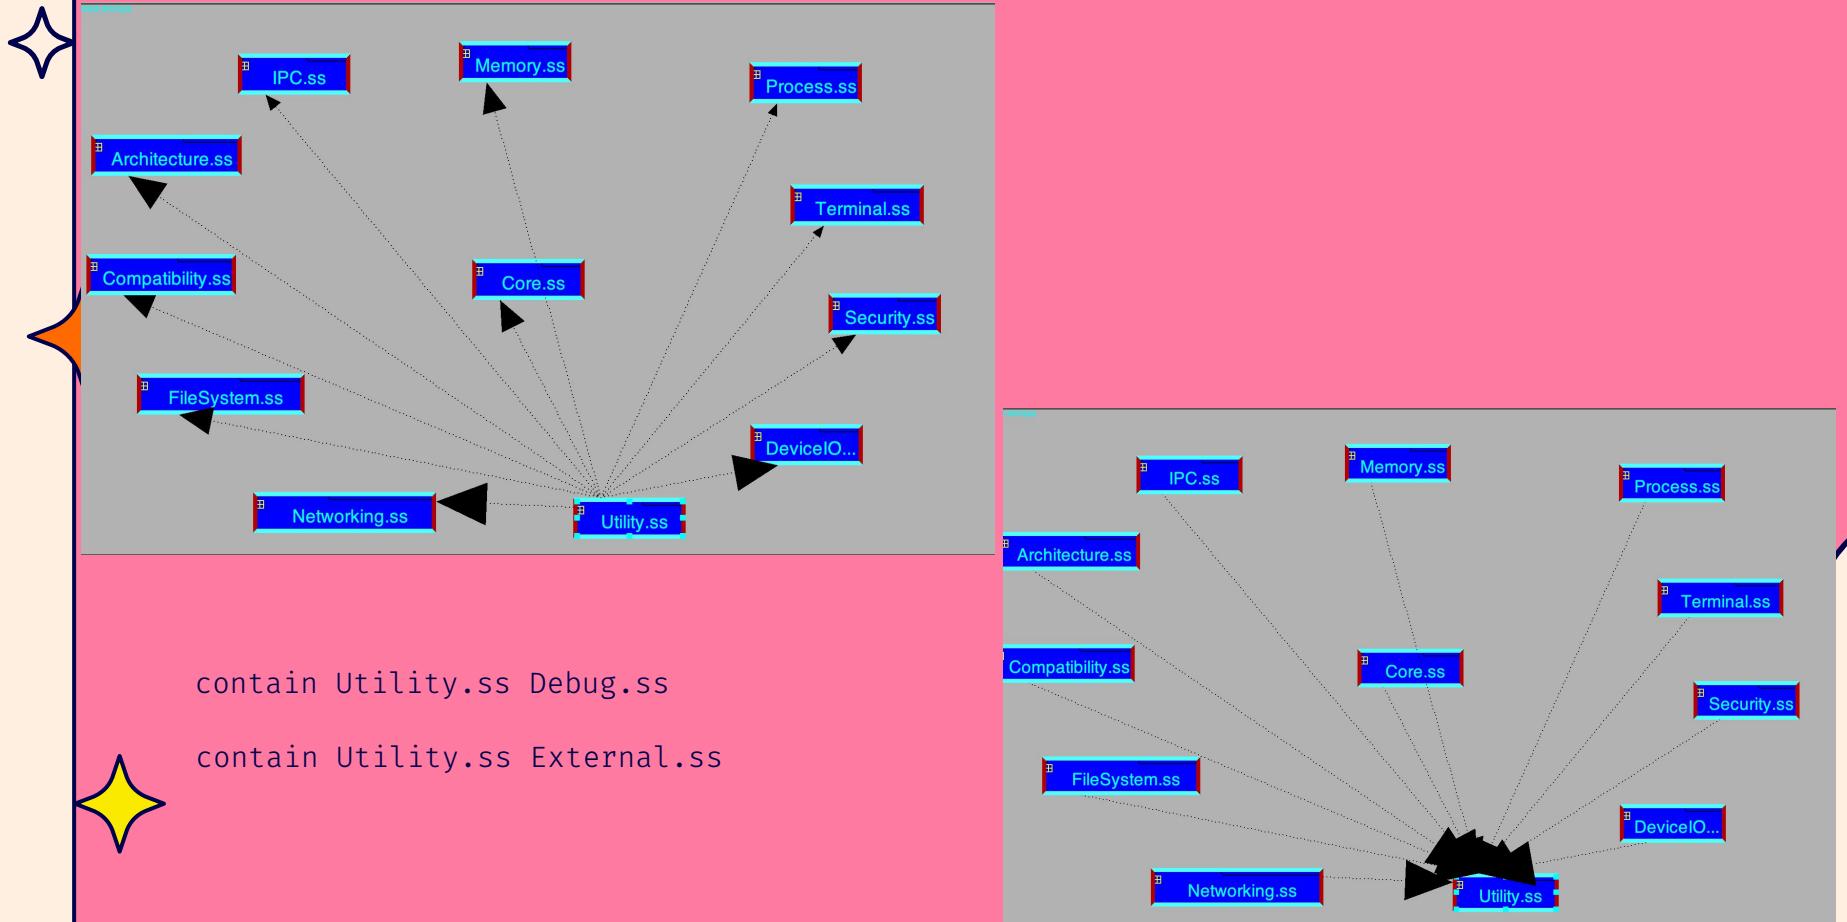


Separation of concern to support other binaries, such as ones built for embedded systems, in the future as well

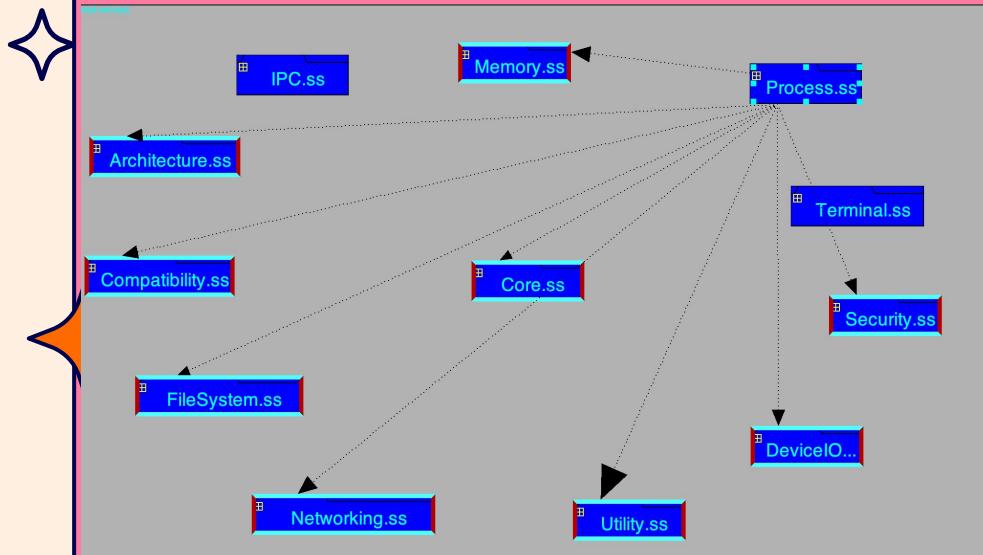
Supporting binaries built for Linux



# UTILITY

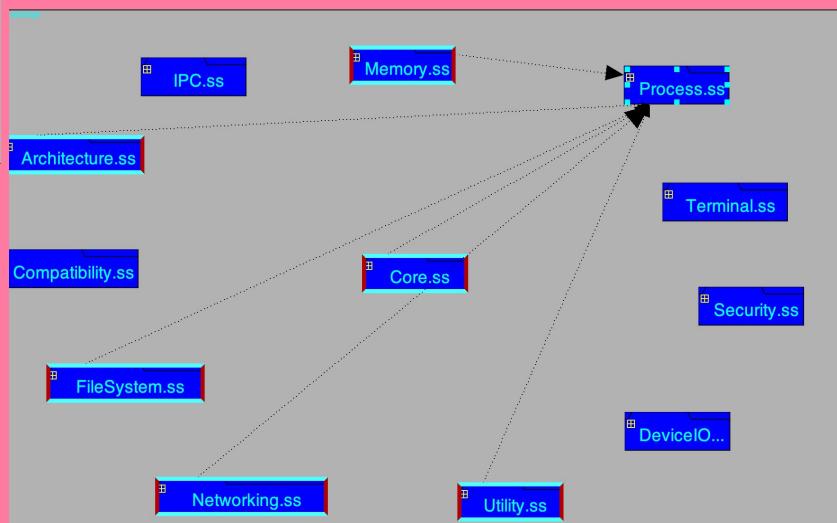


# PROCESS

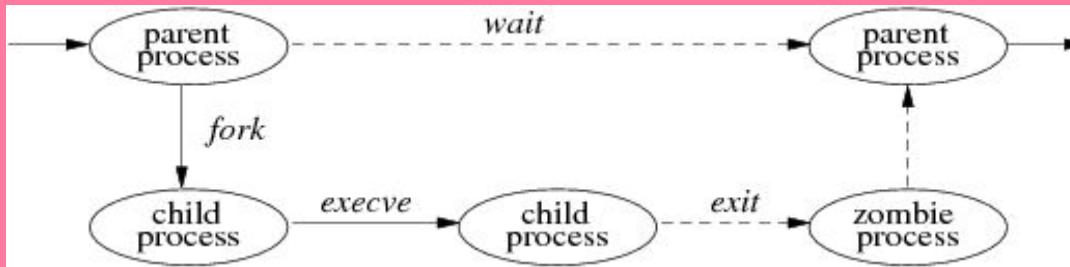


Dependencies on Processes

Dependencies of Processes

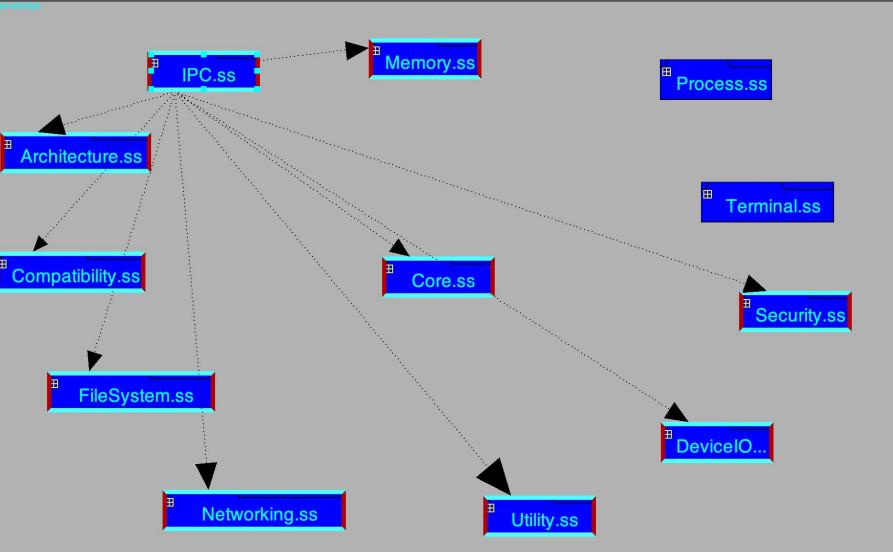


# Process Management



- Processes managed by the kernel
  - Using a priority-based algorithm
- *Jail* system for isolation and accessibility control for different processes
- The image shown depicts the life-cycle of a process
- Use of multiple fork calls for process creation
- Parent-child hierarchical structure for effective running

# IPC

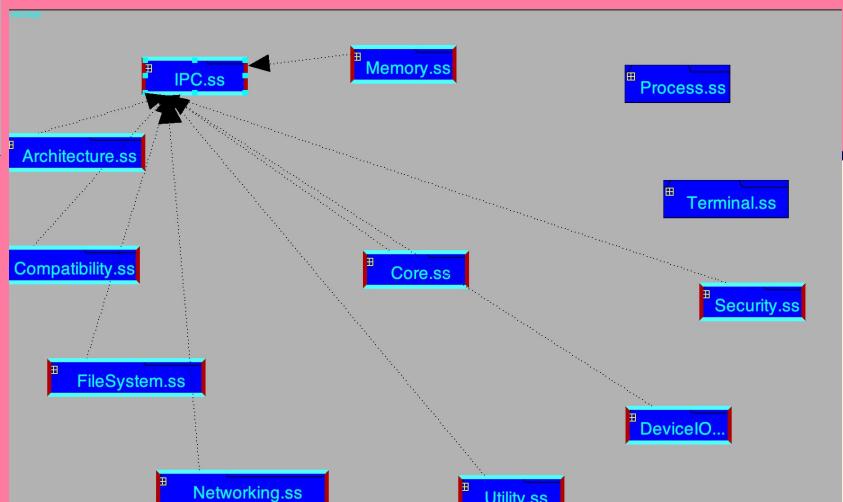


## Protocols

- Sockets must respect the protocols used by the domain

## Sockets

- Abstract data-type used for communicating within the domains



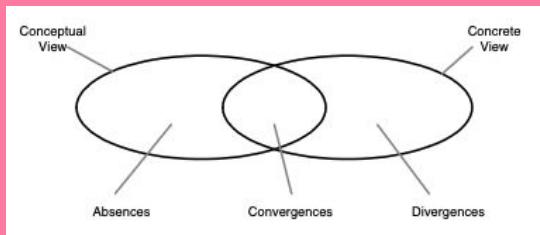
# REFLEXION ANALYSIS

Conceptual  
Model

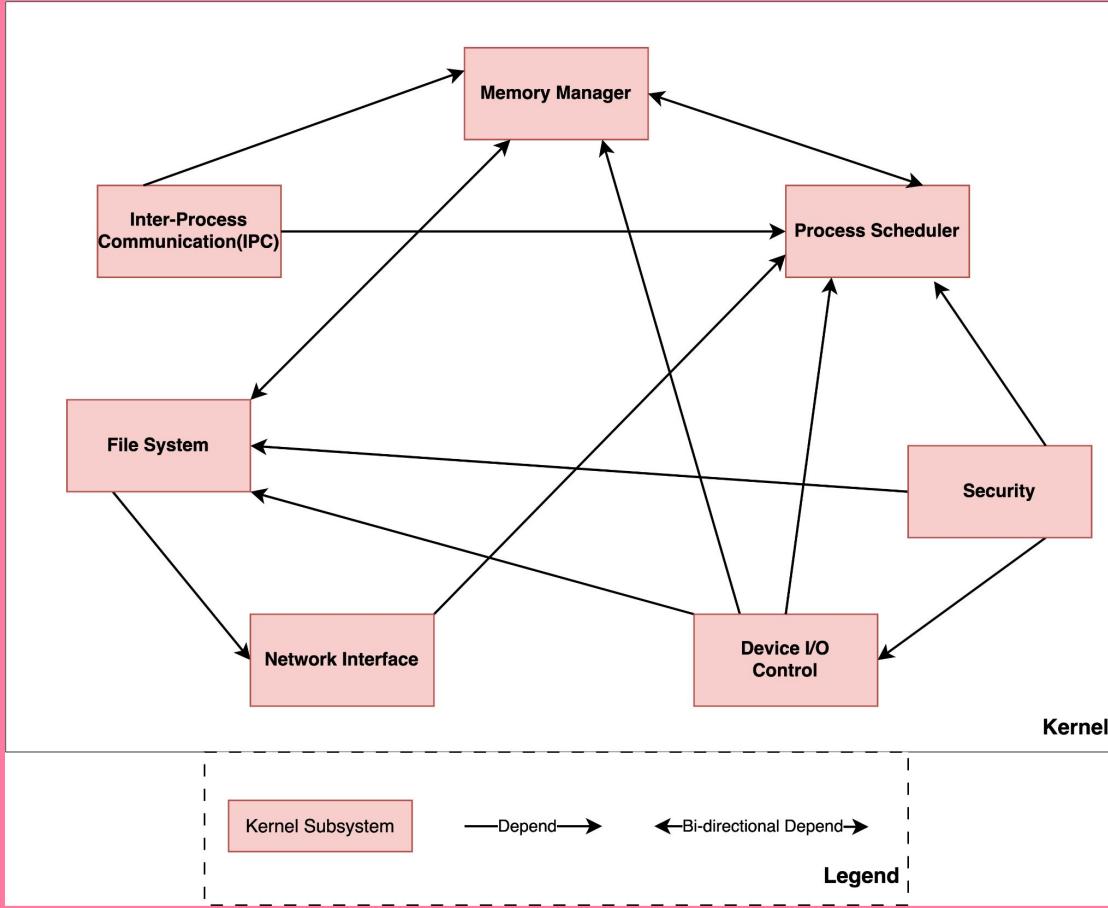
vs

Reflexion  
Model

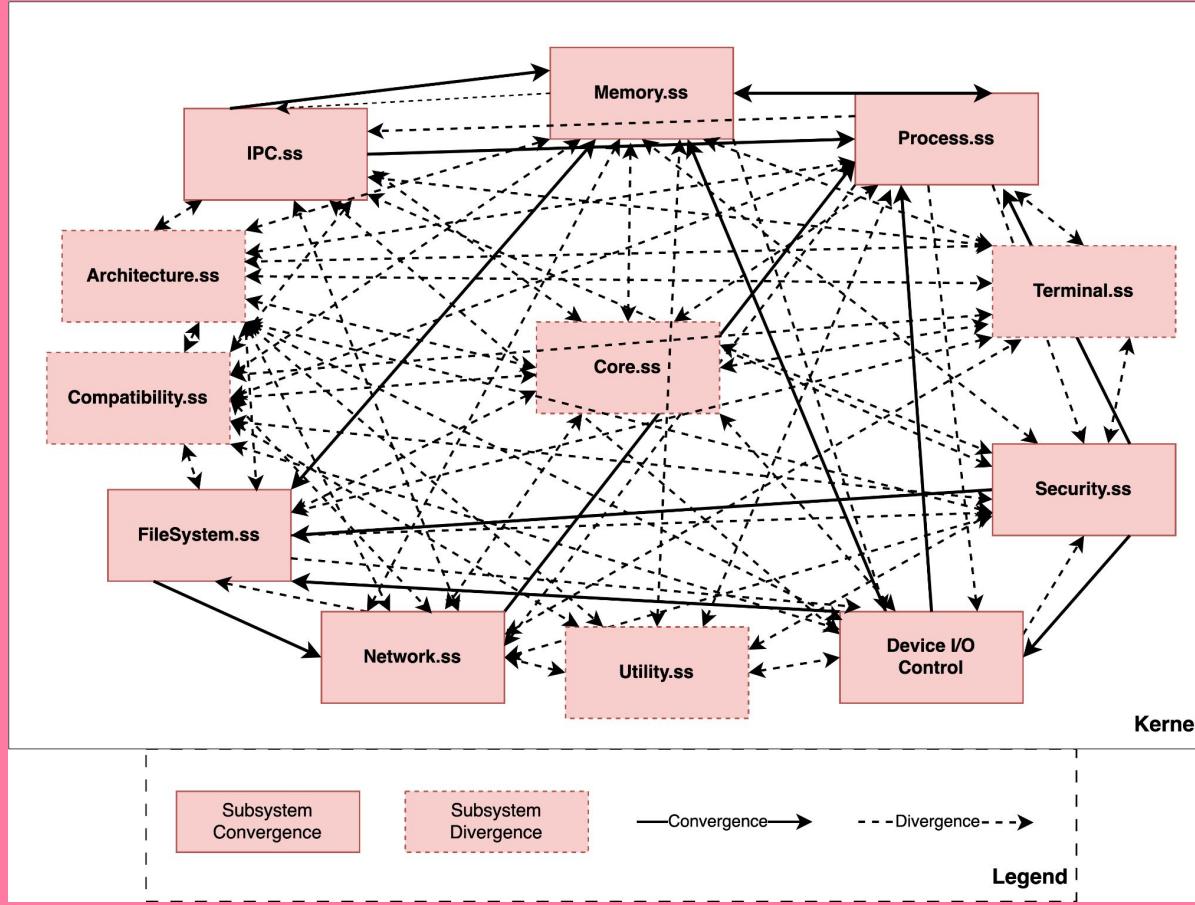
Convergences  
&  
Divergences

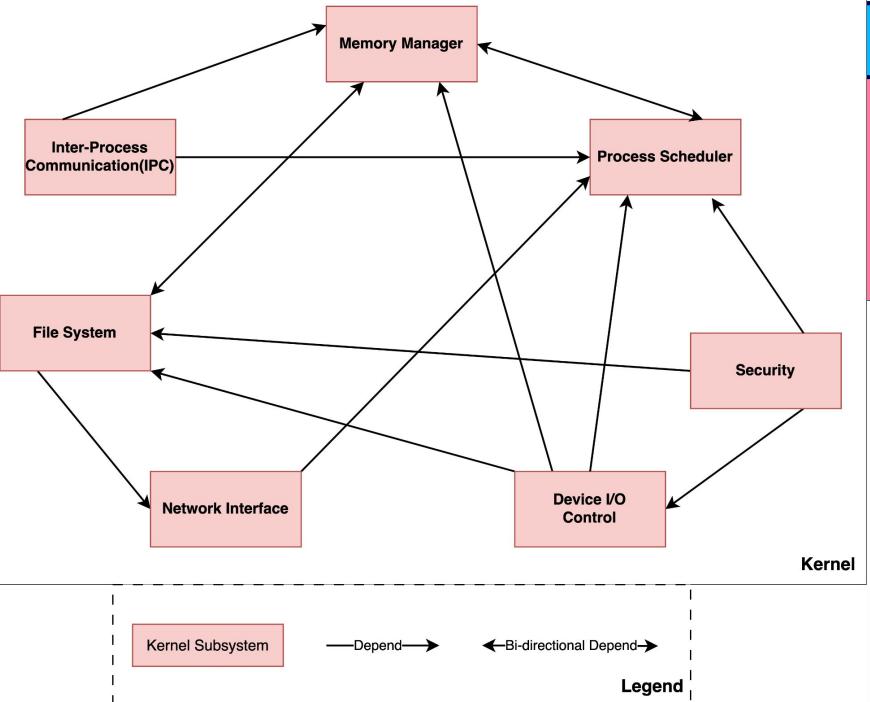


# CONCEPTUAL VIEW

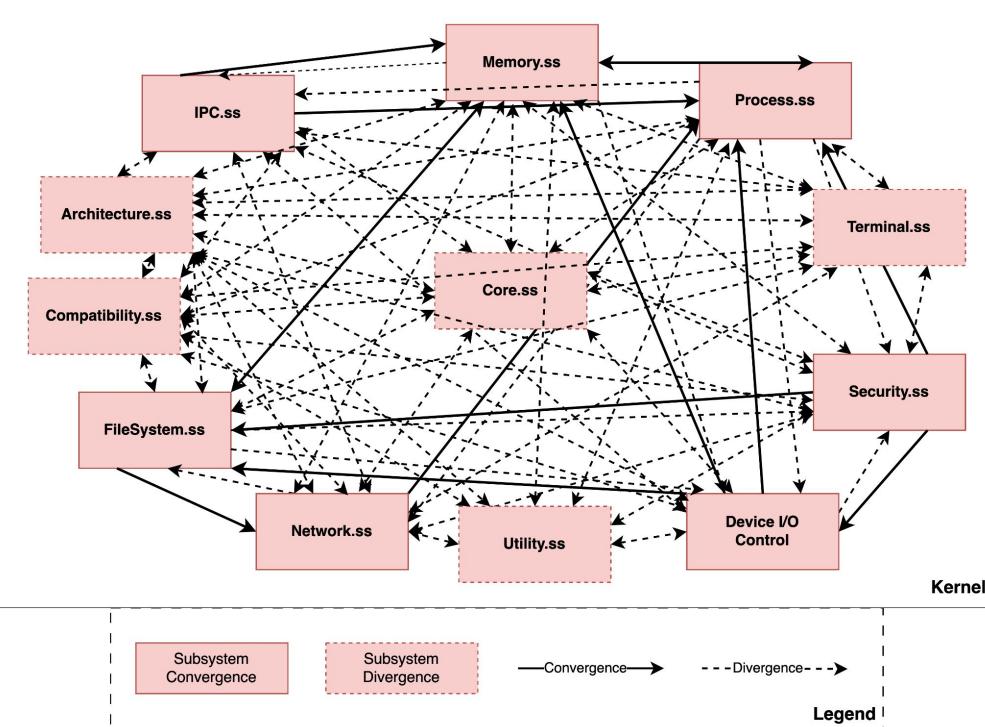


# REFLEXION DIAGRAM





- Always expect more dependencies in concrete architecture
- More functionalities





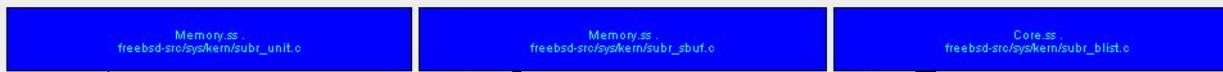
# Concurrency

- In order to make sure that errors don't occur when multiple threads are using the same memory area, or different users are accessing the same file.
- The *mutex* and the *lockf* can be used for this situation.
- *Mutex* can arrange CPU to process the thread without conflict and synchronization
- *lockf* can ensure that when executing a command while holding a file lock

FileSystem.ss  
freebsd-src/sys/devs/devs\_int.h

Memory.ss  
freebsd-src/sys/Memory.ss

Utility.ss  
freebsd-src/sys/Utility.ss

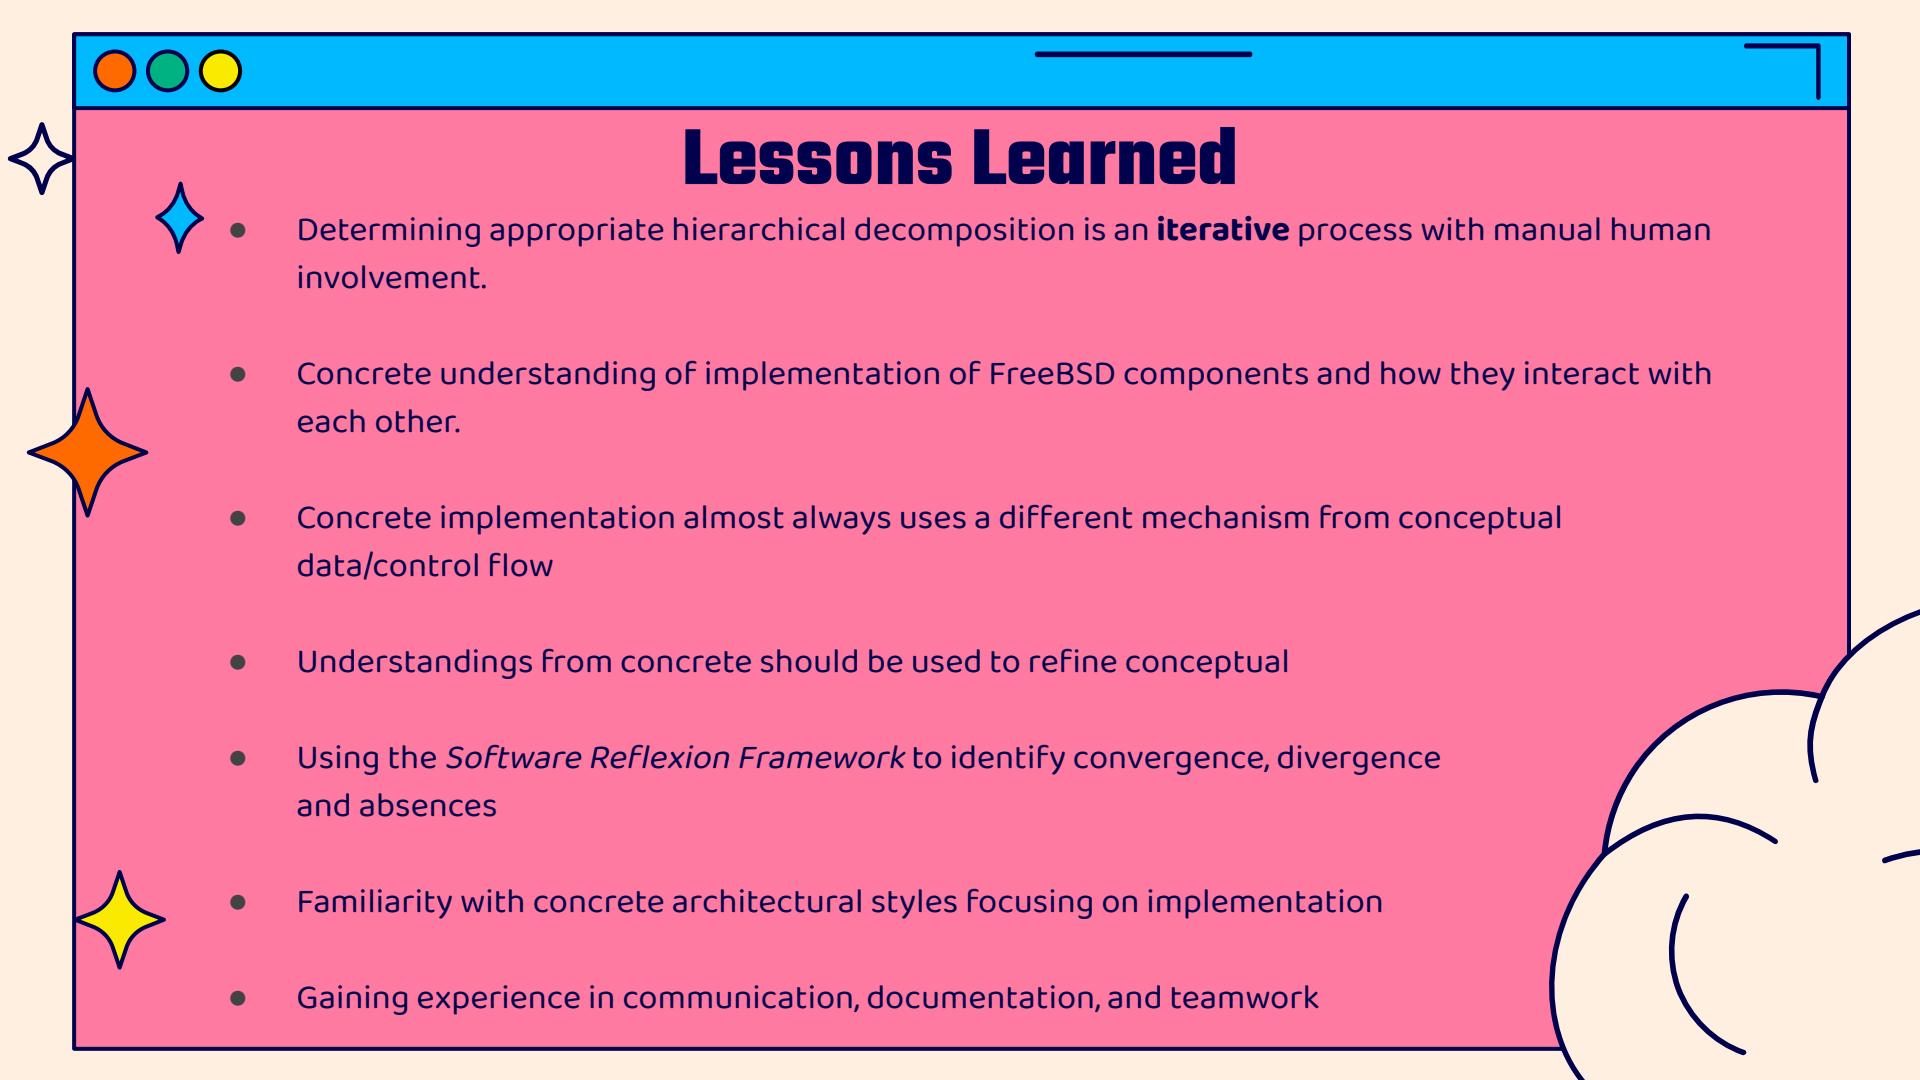


freebsd-src/sys/kern/subr\_unit.c





# LSEdit Demo



# Lessons Learned

- Determining appropriate hierarchical decomposition is an **iterative** process with manual human involvement.
- Concrete understanding of implementation of FreeBSD components and how they interact with each other.
- Concrete implementation almost always uses a different mechanism from conceptual data/control flow
- Understandings from concrete should be used to refine conceptual
- Using the *Software Reflexion Framework* to identify convergence, divergence and absences
- Familiarity with concrete architectural styles focusing on implementation
- Gaining experience in communication, documentation, and teamwork



# Conclusion

- Architecture of the Memory Subsystem and top level subsystems
- Derivation Process
- Reflexion Model
- LSEdit



# References

1. <https://cgit.freebsd.org/src/tree>
2. <https://docs.freebsd.org/en/books/arch-handbook/book/>
3. *Marshall Kirk McKusick, George V. Neville-Neil, Robert N.M. Watson* (2015)  
The Design and Implementation of the FreeBSD® Operating System, 2 e.d.
4. *William Joy, Robert Fabry, Samuel Leffler, M. Kirk McKusick, Michael Karels*,  
Berkeley Software Architecture Manual 4.4BSD Edition  
<https://docs.freebsd.org/44doc/psd/05.sysman/paper.html>
5. *David Garlan, Mary Shaw* (1994) An Introduction to Software Architecture

